

# FTML Project

nelson.vicel-farah, antoine.zellmeyer

June 2022

## 1 Bayes estimator and Bayes risk

### Question 1

- Input space  $\mathcal{X} = [0; 1]$
- Output space  $\mathcal{Y} = \mathbb{R}^+$
- $X$  uniform continuous distribution on  $\mathcal{X}$
- $l(x, y) = \text{squared loss}$
- $Y \sim \text{Exp}(1 + X)$

The Bayes estimator in respect to the squared loss is  $f^*(x) = E[Y|X = x]$ , so

$$f^*(x) = E[Y \sim \text{Exp}(1 + x)]$$

We know that  $E[X \sim \text{Exp}(\lambda)] = \frac{1}{\lambda}$  so

$$\mathbf{f}^*(\mathbf{x}) = \frac{\mathbf{1}}{\mathbf{1} + \mathbf{x}}$$

And the Bayes Risk is

$$\begin{aligned} R^* &= E[l(Y, f^*(X))] \\ &= E_X[E_Y[(Y - f^*(X))^2|X]] \\ &= E_X[\text{Var}(Y|X)] \end{aligned}$$

We know that  $\text{Var}[X \sim \text{Exp}(\lambda)] = \frac{1}{\lambda^2}$

$$\begin{aligned} &= E_X\left[\frac{1}{(1 + X)^2}\right] \\ &= \int_0^1 \frac{1}{(1 + x)^2} dx = \left[-\frac{1}{1 + x}\right]_0^1 \\ &= \frac{1}{2} \end{aligned}$$

## Question 2

Let's define  $\tilde{f}$  as the OLS estimator, which means

$$\tilde{f} = \begin{cases} \mathcal{X} \rightarrow \mathcal{Y} \\ x \rightarrow \hat{\theta}_1 x + \hat{\theta}_2 \end{cases}$$

With  $\hat{\theta}_1$  and  $\hat{\theta}_2$  the scalar parameters that minimize the squared loss. We deduce that

$$\begin{aligned} \hat{\theta}_1 &= \frac{Cov(X, Y)}{Var(X)} \\ &= 12 * Cov(X, Y) \\ &= 12 (E[XY] - E[X]E[Y]) \\ &= 12 (E[E[XY|X]] - E[X]E[E[Y|X]]) \\ &= 12 (E[\mathbf{X}\mathbf{E}[\mathbf{Y}|\mathbf{X}]] - E[X]E[E[Y|X]]) \\ &= 12 \left( E\left[\frac{X}{1+X}\right] - E[X]E\left[\frac{1}{1+X}\right] \right) \\ &= 12 \left( 1 - \log(2) - \frac{1}{2}\log(2) \right) \\ &= \mathbf{12 - 13\log(2)} \end{aligned}$$

And

$$\begin{aligned} \hat{\theta}_2 &= E[Y] - \hat{\theta}_1 E[X] \\ &= E[E[Y|X]] - \hat{\theta}_1 \frac{1}{2} \\ &= \log(2) - \frac{1}{2}(12 - 13\log(2)) \\ &= \frac{\mathbf{15}}{\mathbf{2}}\log(2) - \mathbf{6} \end{aligned}$$

We conclude that

$$\tilde{f}(x) = (12 - 13\log(2))x + \frac{15}{2}\log(2) - 6$$

The simulation in `part1_simulation.py` with the above settings and 10 000 samples of (X,Y) suggests that the Bayes estimator is better at estimating the setting than the OLS estimator. Probably because our model is not linear. Furthermore, the computed generalization error of the Bayes estimator converges to the value of the Bayes Risk we found previously ( $\sim 1/2$ ) which lets us think that the simulation is a good estimation of the theoretical setup.

## 2 Bayes risk with absolute loss

### Question 1

$P(Y|X=x)$  where  $P$  corresponds to an  $\text{Exp}(\lambda)$  continuous distribution.  
The Bayes estimator for  $l_2$  squared loss

$$f_2^*(x) = E[Y|X = x] = E[\lambda e^{-\lambda}] = \frac{1}{\lambda}$$

The Bayes estimator for  $l_1$  absolute loss is the median of  $Y|X=x$  as seen in Question 2

$$f_1^*(x) = \frac{\ln(2)}{\lambda}$$

These Bayes estimators are not equal.

### Question 2

We note  $p = p_{Y|X=x}$

$$\begin{aligned} g(z) &= \int_{\mathbb{R}} |y - z| p(y) dy \\ &= \int_z^{+\infty} (y - z) p(y) dy + \int_{-\infty}^z (z - y) p(y) dy \\ &= \int_z^{+\infty} y p(y) dy - z \int_z^{+\infty} p(y) dy + z \int_{-\infty}^z p(y) dy - \int_{-\infty}^z y p(y) dy \\ \frac{d}{dz} g(z) &= -z p(z) - \left( \int_z^{+\infty} p(y) dy - z p(z) \right) + \left( \int_{-\infty}^z p(y) dy + z p(z) \right) - z p(z) \\ &= \int_{-\infty}^z p(y) dy - \int_z^{+\infty} p(y) dy \end{aligned}$$

Thus,  $\frac{d}{dz} g(z) = 0$  if :

$$\int_{-\infty}^z p(y) dy = \int_z^{+\infty} p(y) dy$$

This occurs when both sides are equal to  $\frac{1}{2}$ , which means that the Bayes estimator  $f^*(x)$  is the median. But to confirm this minimizes  $g(z)$ , let's check the second derivative:

$$\frac{d^2}{dz^2} g(z) = 2p(z) > 0$$

The second derivative is always positive, so our solution is a minimum.



**Step 3 :**

$$\begin{aligned}
E_\epsilon \left[ \frac{1}{n} \|A\epsilon\|^2 \right] &= E_\epsilon \left[ \frac{1}{n} \sum_{i=1}^n (A\epsilon)_i^2 \right] \\
&= E_\epsilon \left[ \frac{1}{n} \text{tr}((A\epsilon)^T (A\epsilon)) \right] \\
&= E_\epsilon \left[ \frac{1}{n} \text{tr}(\epsilon^T A^T A \epsilon) \right]
\end{aligned}$$

The trace is invariant under cyclic permutation.

$$= E_\epsilon \left[ \frac{1}{n} \text{tr}(\epsilon \epsilon^T A^T A) \right]$$

By linearity of the trace and the expected value, we can push the expectation inside.

$$\begin{aligned}
&= \frac{1}{n} \text{tr}(E_\epsilon[\epsilon \epsilon^T] A^T A) \\
&= \frac{1}{n} \text{tr}(\sigma^2 I_n A^T A) \\
&= \frac{1}{n} \text{tr}(A^T A) \sigma^2
\end{aligned}$$

**Step 4 :**

$$\begin{aligned}
A &= I_n - X(X^T X)^{-1} X^T \\
A^T A &= (I_n - X(X^T X)^{-1} X^T)^T (I_n - X(X^T X)^{-1} X^T) \\
&= (\mathbf{I}_n^T - (\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T))^T (I_n - X(X^T X)^{-1} X^T) \\
&= (\mathbf{I}_n - (\mathbf{X}^T)^T ((\mathbf{X}^T \mathbf{X})^{-1})^T \mathbf{X}^T) (I_n - X(X^T X)^{-1} X^T) \\
&= (\mathbf{I}_n - \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T) (I_n - X(X^T X)^{-1} X^T) \\
&= (I_n - X(X^T X)^{-1} X^T)^2 \\
&= I_n - 2X(X^T X)^{-1} X^T + (X(X^T X)^{-1} X^T)^2
\end{aligned}$$



## Step 7

The simulation is in `part3_step7.py`. We suppose :

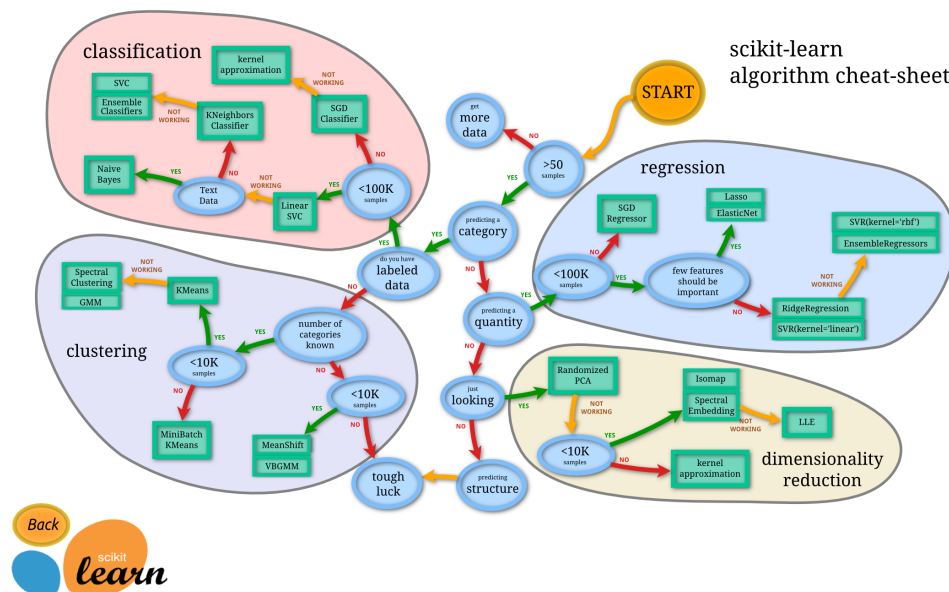
- $n = 10000$
- $d = 20$
- $y = X\theta^*$  with  $\theta^*$  a random vector

We thus estimate  $\sigma^2$  with the result of Step 6. With this setting, the simulation estimates a  $\sigma^2$  of 0, which is an expected result knowing that  $y$  is a simple linear function of  $X$  without any noise added to it (so the variance of  $\epsilon$  is 0).

In a second setting, we define  $y$  as  $y = X\theta^* + \epsilon$  with  $\epsilon$  being a gaussian vector with a standard deviation of 2 and a mean of 0, for which the simulation estimates  $\sigma^2$  as 4, thus being still consistent with the theoretical values.

## 4 Regression

For this section and the next one, we chose `scikit-learn` as our machine learning library because it contains many regression and classification models, it provides easy access to cross validation algorithms and allows us to follow this decision tree :



We should then retrieve the important properties of the dataset :

Number of entries : 1000

Number of features : 20

This needs a regression model that can handle a small set of samples and 20 features

| Model                       | Cross validation score |
|-----------------------------|------------------------|
| SVR Linear                  | 0.8920                 |
| SVR rbf                     | 0.6820                 |
| Ridge alpha=12.0            | <b>0.8932</b>          |
| Ridge alpha=15.0            | <b>0.8932</b>          |
| Lasso alpha=12.0            | 0.8371                 |
| RandomForestRegressor       | 0.8409                 |
| AdaBoost Regressor          | 0.8334                 |
| Gradient Boosting Regressor | 0.8690                 |

Although we should note that the cross validation score is slightly influenced by the random initialization of the models parameters, the Ridge regression seems to be globally the best regression model in this case.

But to find the best alpha parameter of the Ridge model, we can test different values and evaluate them via their cross validation score usgin the `GridSearchCV` from the scikit-learn modules, from which we can deduce that the best alpha parameter is 60.

## 5 Classification

As we did in the previous section, we retrieve useful information about the dataset

Number of entries : 1000

Number of features : 20

This needs a classification model that can handle a small set of samples and 20 features. Let's try some of them:

| Model                      | Cross validation score |
|----------------------------|------------------------|
| LinearSVC                  | <b>0.886</b>           |
| KNeighbors                 | 0.856                  |
| SVC                        | 0.885                  |
| RandomForestClassifier     | 0.858                  |
| AdaBoostClassifier         | 0.864                  |
| GradientBoostingClassifier | 0.858                  |

Although we should note that the accuracy is slightly influenced by the random initialization of models parameters, the linear SVC seems to be globally the best



classification model in this case.

We can finally find the best hyperparameters for **LinearSVC** in the same way as in the previous section. This lets us conclude that the best hyperparameters are : **C=1**, **loss=hinge**, **penalty=L2** as they get a cross validation score of 0.887 (+ 0.1).